

WENHUI ZHU

wzhu59@asu.edu ◊ 602-773-9980 ◊ [GitHub](#) ◊ [Google Scholar](#) ◊ [LinkedIn](#) ◊ [Homepage](#)

RESEARCH INTERESTS

Post-training and alignment for LLMs/MLLMs; agentic RL and Spatial reasoning reasoning; efficient LLMs/MLLMs and Code Agent systems; large-scale multimodal data curation for Foundation Model/MLLMs Scaling.

EDUCATION

Arizona State University *Aug. 2022 – May 2026*
Ph.D. Candidate in Computer Engineering, GPA: 3.89/4.0 Tempe, AZ

Arizona State University *Aug. 2020 – May 2022*
M.S. in Computer Engineering Tempe, AZ

Northern Arizona University & CQUPT *Aug. 2014 – May 2019*
B.S. in Electrical Engineering Flagstaff, AZ & Chongqing, China

RESEARCH AND INDUSTRY EXPERIENCE

LinkedIn Corporation *Jan. 2026 – Present*
Senior AI Scientist Mountain View, CA

- Optimize GRPO pipelines that accelerate time-to-peak training by 57% while improving peak model quality by an additional 2.7% relative over the previous post-training pipeline for LinkedIn Recruiter Hiring LLM models.
- Proposed efficient long-reasoning inference, a training-free payload-aware KV-cache eviction policy that improves real-world agentic usage accuracy from 76.6% to 81.4% at a 512 active-KV budget and reduces max-length hits from 78 to 55 compared with R-KV.
- Develop efficient black-box distillation methods, which reduces per-step wall-clock time by up to $8\times$ over vanilla GAD while improving quality by up to 3% on Qwen2.5-7B.
- Contributed to RecLLM, LinkedIn recommendation-context reasoning LLM for behavior information and SIDs for recommendation understanding, offline judging, and RL feedback; implemented the text-SID alignment training pipeline.

LinkedIn Corporation *May 2025 – Aug. 2025*
Generative AI Intern Sunnyvale, CA

- Led EVTP-IVS, a visual token pruning method for instruction visual segmentation MLLMs that uses only 20% of visual tokens while achieving up to $5\times$ speed-up on video tasks and $3.5\times$ on image tasks with comparable accuracy. Adapted this to downstream data-curation pipelines, reducing the cost of scaling multimodal datasets.

Ira A. Fulton Schools of Engineering, Arizona State University *May 2022 – Dec. 2025*
Research Assistant, Geometry Systems Laboratory (GSL Lab) Tempe, AZ

- Built domain multimodal foundation-model systems spanning VLM post-training, large-scale data curation, preference assistants, agent-led RAG, generative modeling, and rigorous evaluation.
- Published first/co-first-author work at ICCV, ECCV, ICML, CVPR, AAAI, COLM and ACL, with broader collaborations across MICCAI, IPMI, and JAMA Ophthalmology.

SELECTED PUBLICATIONS (*: EQUAL CONTRIBUTION)

Post-Training and Alignment

- **“S-SPPO: Semantic-Calibrated Self-Play Preference Optimization”**
X. Chen*, **W. Zhu***, J. Wang, P. Qiu, Z. Wang, H. Li, Z. He, X. Dong, P. Tiwari, M. Xu, Y. Xiong, F. Luo, A. Razi, B. H. Rappazzo, A. Schneider, et al.
ICML 2026
- **“DRA-GRPO: Exploring Diversity-Aware Reward Adjustment for R1-Zero-Like Training of Large Language Models”**
X. Chen*, **W. Zhu***, P. Qiu, X. Dong, H. Wang, H. Wu, H. Li, A. Sotiras, Y. Wang, et al.
ACL 2026
- **“Sampling for Quality: Training-Free Reward-Guided LLM Decoding via Sequential Monte Carlo”**
J. Markovic-Voronov*, **W. Zhu***, B. Long, Z. Wang, S. S. Gu, K. Behdin, et al.
COLM 2026

- **“Prompt-OT: An Optimal Transport Regularization Paradigm for Knowledge Preservation in Vision-Language Model Adaptation”**
X. Chen*, **W. Zhu***, P. Qiu, H. Wang, H. Li, H. Wu, X. Dong, A. Sotiras, Y. Wang, et al.
WACV 2026
- **“AHA: Aligning Large Audio-Language Models for Reasoning Hallucinations via Counterfactual Hard Negatives”**
Y. Chen, **W. Zhu***, X. Chen, Z. Wang, X. Li, P. Qiu, H. Wang, X. Dong*, Y. Xiong, A. Schneider, Y. Nevmyvaka, & Y. Wang
ACL 2026

Agentic Reinforcement Learning and Systems

- **“AriadneMem: Threading the Maze of Lifelong Memory for LLM Agents”**
W. Zhu, X. Chen, Z. Wang, J. Wang, X. Dong, M. Huang, R. Cai, H. Sang, et al.
arXiv preprint, 2026
- **“Your Agent is More Brittle Than You Think: Uncovering Indirect Injection Vulnerabilities in Agentic LLMs”**
W. Zhu, X. Dong, X. Chen, R. Cai, P. Qiu, Z. Wang, O. Frunza, S. Tang, J. Gu, et al.
arXiv preprint, 2026
- **“Mags-RL: Wearing Multimodal LLMs a Magnifying Glass via Agentic Reinforcement Learning for Complex Scene Reasoning”**
X. Dong*, **W. Zhu***, P. Qiu, X. Chen, X. Yu, X. Li, Z. Wang, S. Tang, G. Li, Y. Xiong, H. Wang, Y. Chen, P. Tiwari, & Y. Wang
arXiv preprint, 2026
- **“Context Pruning for Coding Agents via Multi-Rubric Latent Reasoning”**
J. Wang*, X. Chen*, **W. Zhu***, H. Li, Z. Wang, Z. He, F. Cai, A. S. Carreon-Rascon, X. Dong, et al.
arXiv preprint, 2026
- **“EZBlender: Efficient 3D Editing with Plan-and-ReAct Agent”**
H. Wang*, **W. Zhu***, S. Tang, Z. Wang, X. Dong, X. Li, X. Chen, A. Bastola, et al.
WACV 2026 Oral

Efficient MLLMs/LLMs

- **“CE-GAD: Curriculum-Efficient On-Policy Adversarial Distillation of Large Language Models”**
W. Zhu, H. Sang, X. Chen, S. M. Hamidi, H. Yu, Z. Wang, J. Wang, Y. Xu, X. Dong, Z. Zhou, R. He, J. Markovic-Voronov, K. Behdin, S. Ghosh, & A. Geramifard
2026
- **“SODA: Semi On-Policy Black-Box Distillation for Large Language Models”**
X. Chen*, J. Wang*, **W. Zhu***, P. Qiu, X. Dong, H. Sang, Z. Wang, A. Geramifard, & F. Luo
2026
- **“OTPrune: Distribution-Aligned Visual Token Pruning via Optimal Transport”**
X. Chen*, **W. Zhu***, G. Li, X. Dong, Y. Xiong, H. Wang, P. Qiu, Q. Song, Z. Wang, et al.
WACV 2026
- **“EVTP-IVS: Effective Visual Token Pruning for Unifying Instruction Visual Segmentation in Multimodal Large Language Models”**
W. Zhu, X. Chen, Z. Wang, S. Tang, S. Ghosh, X. Dong, R. Koner, & Y. Wang
WACV 2026

MLLMs/Foundation Model and Large-Scale Data Curation Scaling

- **“Ophth-500K: Curating Unstructured Video Streams for Scaling Multimodal Ophthalmology Foundation Models”**
H. Wang, X. Li, **W. Zhu***, X. Dong, X. Chen, Y. Xiong, L. Liu, A. Razi, O. Dumitrascu, & Y. Wang
2026
- **“OphIn-500K: Curating Web-Scale Visual Instructions for Scaling Ophthalmic Multimodal Large Language Models”**
X. Dong, **W. Zhu***, X. Chen, H. Wang, X. Li, Y. Xiong, J. Cheng, J. Wang, X. Yu, H. Wu, S. Tang, Z. Wang, L. Liu, S. Lin, O. Dumitrascu, & Y. Wang
2026
- **“LLaDA-MedV: Exploring Large Language Diffusion Models for Biomedical Image Understanding”**
X. Dong, **W. Zhu***, X. Chen, Z. Wang, P. Qiu, S. Tang, X. Li, & Y. Wang
CVPR 2025

Machine Learning

- **“Cracking Instance Jigsaw Puzzles: An Alternative to Multiple Instance Learning for Whole Slide**

Image Analysis”

X. Chen, P. Qiu, **W. Zhu***, H. Wang, H. Li, X. Dong, X. Sun, X. Yu, Y. Wang, A. Razi, & A. Sotiras
ICCV 2025

- **“FIC-TSC: Learning Time Series Classification with Fisher Information Constraint”**
X. Chen, **W. Zhu***, P. Qiu, H. Wang, H. Li, Z. Li, Y. Wang, A. Sotiras, & A. Razi
ICML 2025
- **“How Effective Can Dropout Be in Multiple Instance Learning?”**
W. Zhu, P. Qiu, X. Chen, Z. Yang, A. Sotiras, A. Razi, & Y. Wang
ICML 2025
- **“DGR-MIL: Exploring Diverse Global Representation in Multiple Instance Learning for Whole Slide Image Classification”**
W. Zhu, X. Chen, P. Qiu, A. Sotiras, A. Razi, & Y. Wang
ECCV 2024
- **“TimeMIL: Advancing Multivariate Time Series Classification via a Time-Aware Multiple Instance Learning”**
X. Chen, P. Qiu, **W. Zhu***, H. Li, H. Wang, A. Sotiras, Y. Wang, & A. Razi
ICML 2024

INVITED TALKS

- Nov.2026: Next-Generation AI & Machine Learning | Boston, MA
- Jan.2026: Oral Presentation, WACV 2026.
- Oct.2025: Award Talk and Challenge Presentation, MICCAI MuCaRD Challenge, MICCAI 2025.
- Oct.2024: Oral Presentation, AAAI 2025.

SELECTED HONORS AND AWARDS

- Outstanding Graduate Student Award, School of Computing and Augmented Intelligence (SCAI), Arizona State University, 2026. Sole recipient per department.
- 1st Place Winner, MICCAI MuCaRD Challenge, Final Task, 2025.
- SCAI Doctoral Fellowship, Arizona State University, Spring 2025.
- Outstanding Contribution Award, MICCAI UWF4DR Challenge, Tasks 1 and 3, 2024.
- Two 3rd Place Awards, MICCAI UWF4DR Challenge, Tasks 1 and 3, 2024. Out of 22 teams.
- Two 3rd Place Awards, MICCAI MMAC Challenge, Tasks 1 and 3, 2023. Out of 59 teams.
- SCAI Doctoral Fellowship, Arizona State University, Spring 2024.
- ASU Graduate Research Fellowship / Scholarship.
- WACV Travel Award 2026 and Travel Awards by Arizona State University.

TECHNICAL STRENGTHS

- **Research areas:** MLLM post-training, reinforcement learning for VLMs, multimodal RAG, agentic visual reasoning, large-scale multimodal data curation, Foundation Model Scaling, efficient MLLMs.
- **Modeling:** transformers, diffusion models, optimal transport, visual token pruning, multiple instance learning, multimodal benchmark design.
- **Engineering:** Python, PyTorch, TensorFlow, OpenCV, Java, C, large-scale model evaluation, data curation, vLLM, Verl and SGLang.

ACADEMIC SERVICES

- Conference Reviewer: ACL, EMNLP, NeurIPS, ICML, ICLR, CVPR, ICCV, ECCV, COLM, MICCAI, AAAI, and WACV.
- Journal Reviewer: JAMA Ophthalmology, IEEE Transactions on Medical Imaging (TMI), IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI), IEEE Transactions on Artificial Intelligence (TAI), Medical Image Analysis (MedIA), and Transactions on Machine Learning Research (TMLR).

TEACHING

- Teaching Assistant, Ira A. Fulton Schools of Engineering, School of Computing and Augmented Intelligence, Arizona State University.